

Types of Data Storage

- SSD Drives
- Spindle Magnetic Drives
 - High Speed - 15,000 RPM spindle speed for high access rate data (600GB)
 - 10,000 RPM spindle speed drives for historical data (1 TB)
 - Common Disk Drives – 7,200 and 5,400 spindle speed drives (1-8 TB)
- Storage Area Network (SAN)

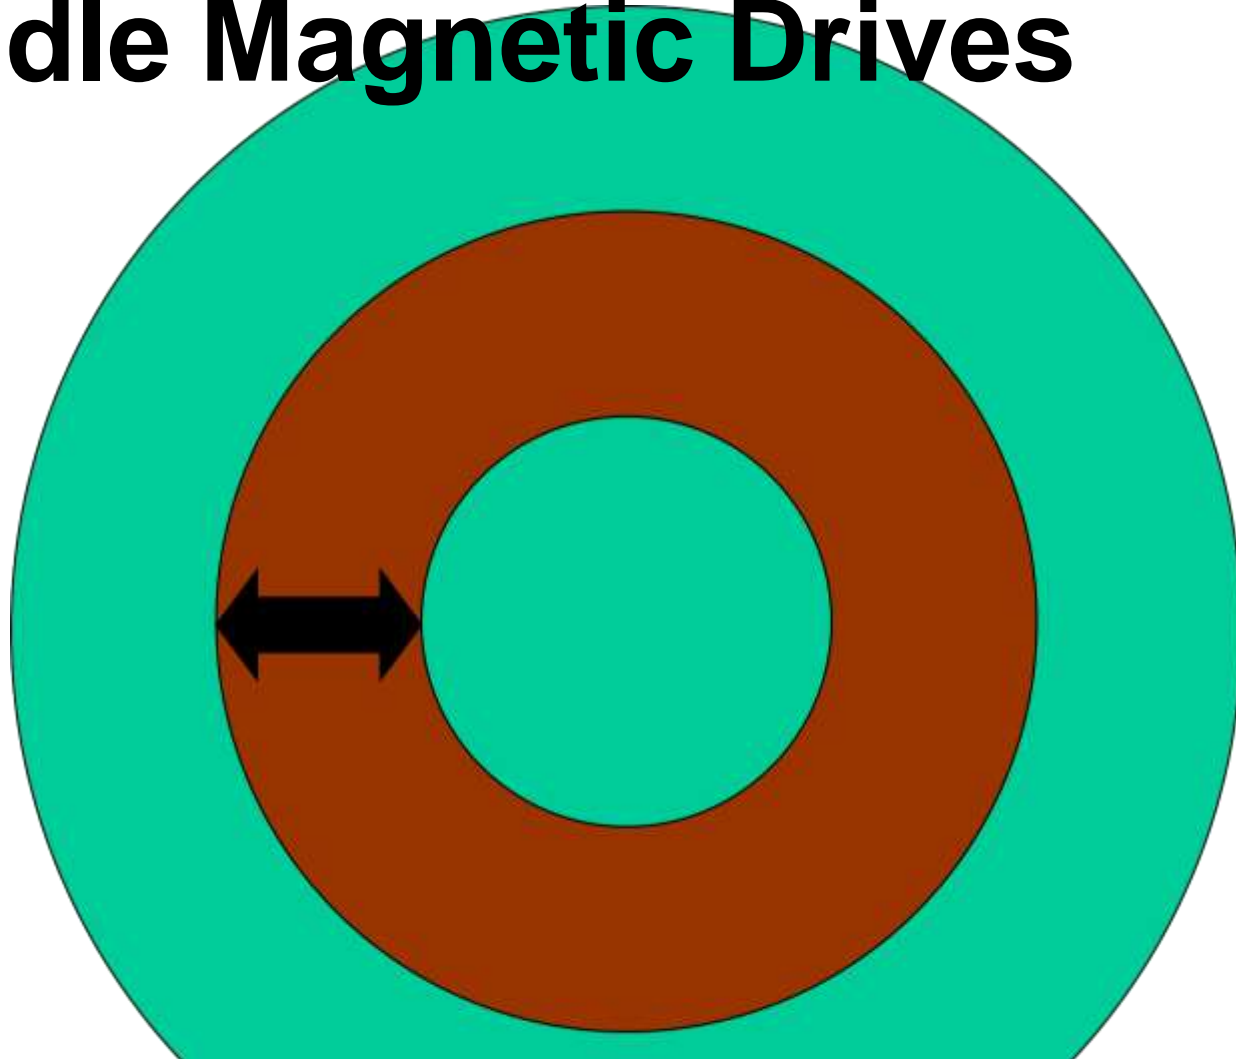
Spindle Magnetic Drives



Disk is organized into sectors.
The disk arm moves to a spot
to read a byte of data

Western Digital VelociRaptor 300 GB,Internal,10000 RPM,3.5" (WD3000BLFS) Hard Drive

Spindle Magnetic Drives



Disk Layout - The FASTEST location on a disk is where the disk arm has to move the least to read or write MOST data

Solid State Disk (SSD)

A **solid-state drive (SSD)** is a nonvolatile storage device that stores persistent data on **solid-state** flash memory. **Solid-state drives** actually aren't **hard drives** in the traditional sense of the term, as there are no moving parts involved



Disk is organized into cells.
Each byte is directly
addressable and readable.

Solid State Disk (SSD)

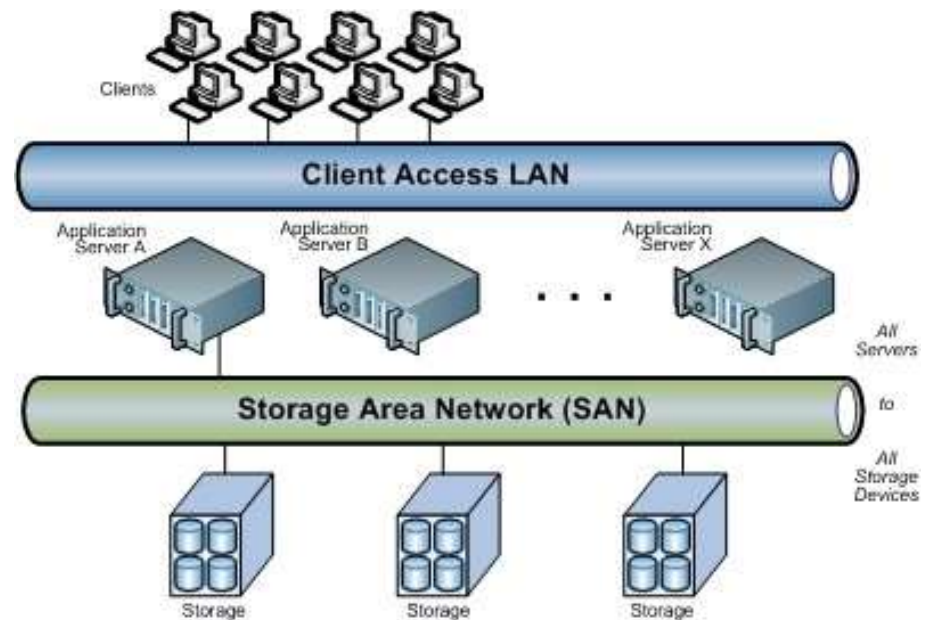
- Performance varies widely by type and manufacture (larger size is faster)
- Drive lifespan and reliability is improving
- Minimum 3 times faster than a hard disk on the same machine
- New SSD drive with M.2 interface will be much faster...

Fusion Drives

- Combine SSD drives as a cache with a slower spindle disk drive
- Most used data is in the cache
- Works great until your active data set exceeds the size of the cache
- Requires CPU cycles to sync cache with disk drive

Storage Area Network (SAN)

- High-speed network that provides block-level access to storage
- Composed of hosts, switches, storage elements, and storage devices that are interconnected

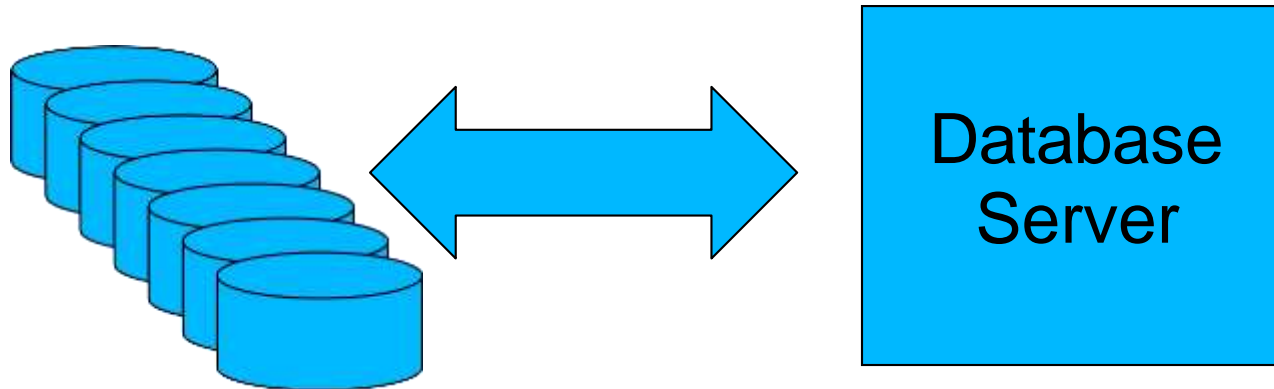


Disk Interface

- SCSI – Small Computer System Interface
- SATA – Serial ATA (3 GB/s – 6 GB/s)
- SAS – Serial Attached SCSI (6 GB/s – 12 GB/s)
- Fiber Channel – High-speed network technology (32-128 GB/s)
- Thunderbolt – Developed by Intel and Apple (10 GB/s – 40 GB/s)

Best Performance

- Multiple Disks attached to Fast Interfaces



What is RAID?

- Redundant Array of Independent Disks
- Goal is to increase:
 - Performance
 - Reliability
 - Safety of Data

RAID 1 - Safest

- Exact copy of two or more disks
- Safe – if one disk fails the other will continue
- Fast – reads can be twice as fast (reads from both disks at the same time)
- The array will continue to operate as long as at least one member drive is operational

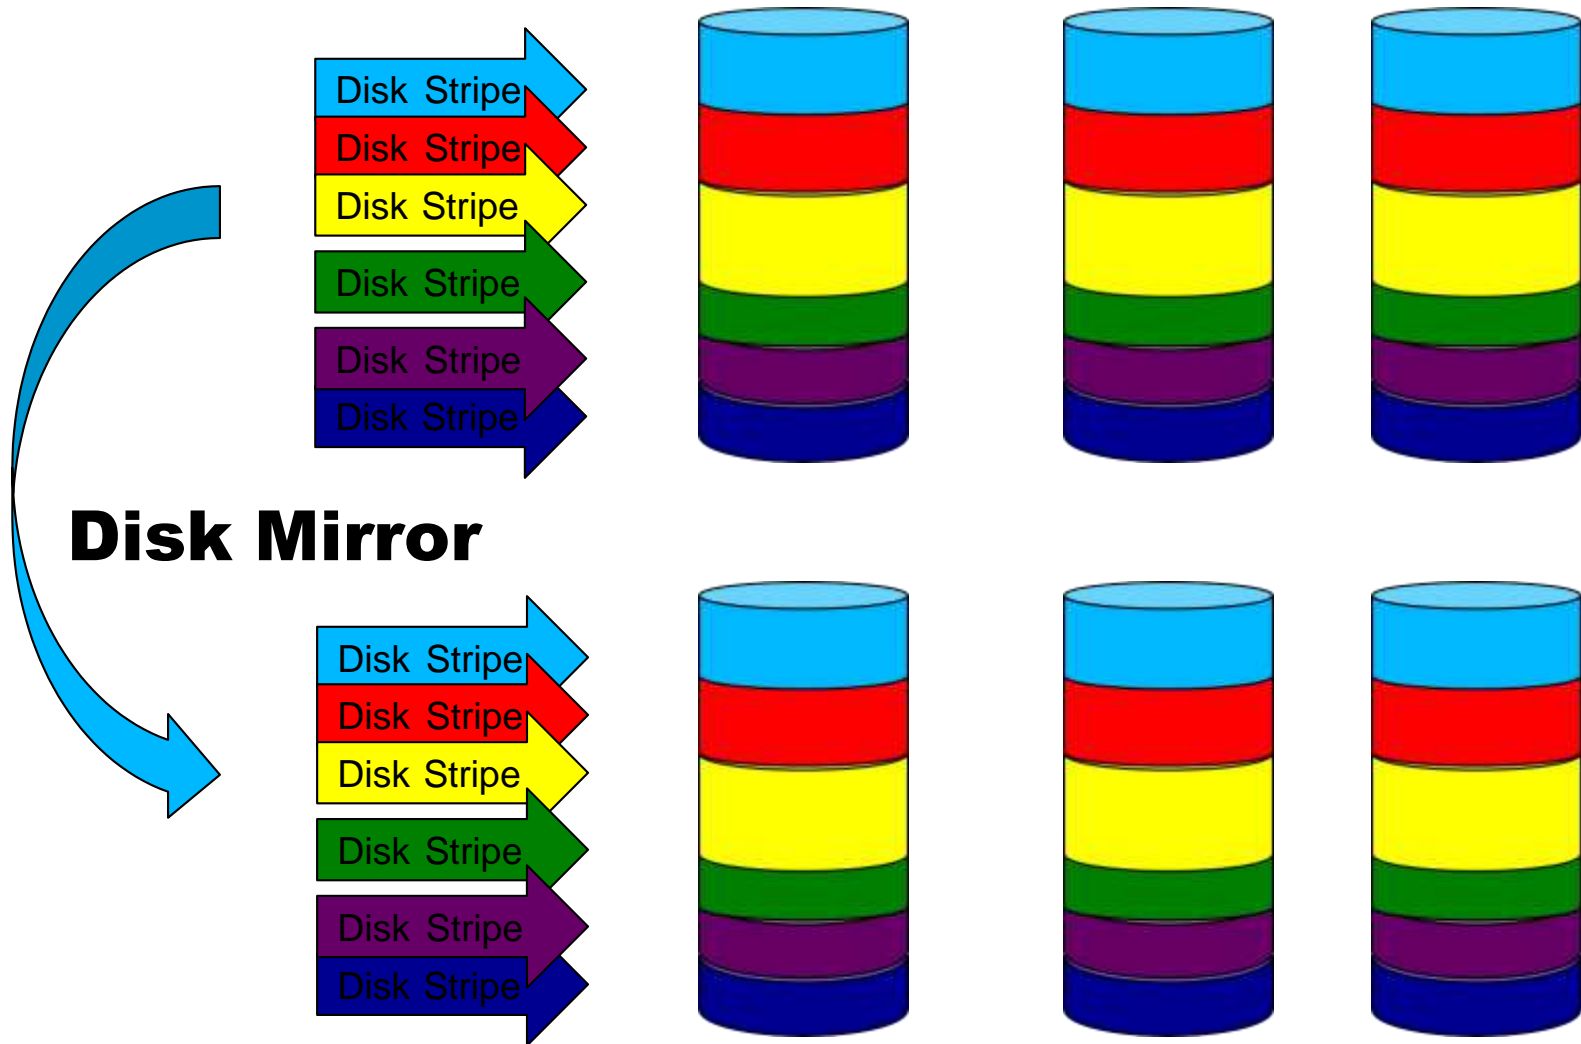
RAID 0 - Fast

- Stripes data evenly across two or more disks
- No redundancy and no fault-tolerance
- Performance – provides data at rates up to 'n' times faster where 'n' is the number of disks

RAID 10 – Safest + Fastest

- Combines the Mirroring of RAID 1 with the Striping of RAID 0
- Best Performance
- Best Reliability and Safety
- Best Configuration is a smaller Block/Strip Size
 - Recommend 64K or 128K

RAID 10 Configuration



RAID 10 Configuration

- Recommended Stripe/Block Size for Databases with a 2K to 16K Page Size
 - 64K best
 - 128 K
 - 256 K is the max Size
- Filesystems benefit from a larger Stripe/Block Size – NOT Databases

RAID 5, 6, 7... etc. - Dangerous

- **RAID 5** is block-level striping with distributed parity. It requires that all drives except one be present to operate.
- Upon failure of a single drive, reads can be calculated from the distributed parity such that no data is lost.
- ***Parity Calculation is slow and may fail***

Do NOT Use RAID 6,5, .. etc

- Why is RAID5 more popular than RAID10 among storage administrators?
 - It requires fewer drives to deliver the same storage capacity.
 - Storage salespeople can present a less expensive proposal to meet the required storage volume. They make a slanted case to make a sale!
 - Most people have never studied the issue themselves and so trust that RAID5 is good.

CERN Disk Drive Failure Rates

- Testing at CERN after they experienced data loss on RAID5 arrays determined:
 - Most drive failures (80%) are caused by hardware and firmware failure (another 10% from wrong firmware version).
 - Partial media failure accounts for much of the rest of the data loss experienced on both magnetic and SSD drives as they aged.
 - They experienced cosmic ray damage flipping bits, equally on both magnetic and SSD type drives.

CERN Disk Drive Failure Rates

- Drives today are commodity priced. Is the risk of data loss worth the relatively small savings?
- The failure rates and expected lifespans for “premium” drives are identical to commodity retail drives. So, it doesn’t help that you are spending >\$1000 per drive.
- According to a storage industry study, failure of a second drive is 4X more likely than the single drive failure rate would predict!
- Atomic writes across multiple drives in a RAID5 array are not guaranteed!

CERN Disk Drive Failure Rates

- Larger drives take longer to rebuild increasing the likelihood of losing a second drive.
- A recent study concluded that drives over 1TB are statistically likely to suffer from unrecoverable multiple bit dropouts.
- The number of bits on the drive exceeds the bit failure rate!
- The error rates as observed by the CERN study on silent corruption are far higher than the official rate of one in every 10^{16} bits.
- The observed error rate was about one in 10^7 bits, or 1 out of about 1 in every 1,000,000 bits (~125,000 bytes).

RAW vs Cooked Space

- RAW – Informix has direct access to the device and space
 - Used to be 25% faster
 - Support for RAW devices is dwindling
 - Harder to manage
- Cooked – Informix accesses the UNIX Filesystem to access a space

RAW Space

- RAW – Informix has direct access to the device
- No UNIX OS Overhead
- No UNIX OS Buffering
- No UNIX Management of devices
- Used to be 25% faster
- Support for RAW devices is dwindling

Using Raw Devices

- When most UNIX systems create a device, they will create two means of accessing that device.
- Block mode of access. This is used by the UNIX file system. The device will have a name that does not begin with the letter “r” and a permissions display in “ls” that begin with the letter “b”. The following is an example:

```
brw-----      1 sysinfo sysinfo  1, 15 Jun 21 1995 /dev/dsk/0s1
```

- ***Raw mode of access. The device will have a name that begins with “r” and the permission displays will begin with the letter “c”. The following is an example:***

```
crw-----      1 sysinfo sysinfo  1, 15 Jun 21 1995 /dev/rdisk/0s1
```

Always use raw mode device for your Informix chunks.

Using Raw Devices

- Create the Raw device with the disk partition utility

Expert Partitioner

System View

- train6
 - Hard Disks
 - + sda
 - sdb
 - sdb1
 - sdb2
 - RAID
 - Volume Management
 - Crypt Files
 - Device Mapper
 - NFS
 - Btrfs
 - tmpfs
 - Unused Devices

Hard Disks

Device	Size	F	Enc	Type	FS Type	Label	Mount Point	Start	
/dev/sda	298.09 GiB			ST3320820AS				0	3
/dev/sda1	100.01 GiB			Linux native	Btrfs		/	0	1
/dev/sda2	8.00 GiB			Linux swap	Swap		swap	13055	1
/dev/sda3	190.08 GiB			Linux native	XFS		/save	14100	3
/dev/sdb	298.09 GiB			ST3320820AS					2
/dev/sdb1	160.00 GiB			Linux LVM					
/dev/sdb2	138.09 GiB			Linux native	Ext2		/informixchunks	20887	3

RAW

Using Raw Devices

- Do NOT Format
- Do NOT Mount

Formatting Options

☐ Format partition
File System
Btrfs ↓
[Options...]
☒ Do not format partition
File system ID:
0x8E Linux LVM ↓
[] Encrypt Device

Mounting Options

☐ Mount partition
Mount Point
↓
[Fstab Options...]
☒ Do not mount partition

[Subvolume Handling]

Using Raw Devices

- Review your OS for specific steps to create Raw Devices

Cooked Space

- Cooked – All Informix access to the device is through the OS
- OS Overhead and OS Buffering
- Simple Management of Devices
- May be Slower
- Pick your filesystem carefully – do not use a Journalled Filesystem

Using Cooked Devices

- Avoid Journaling Filesystems
 - Worst - EXT4, EXT3 (with journaling enabled), ZFS, BTFS
 - Acceptable - JFS2/OpenJFS
- Best - On Linux: EXT2 or EXT3 with journaling disabled

Creating Cooked Files

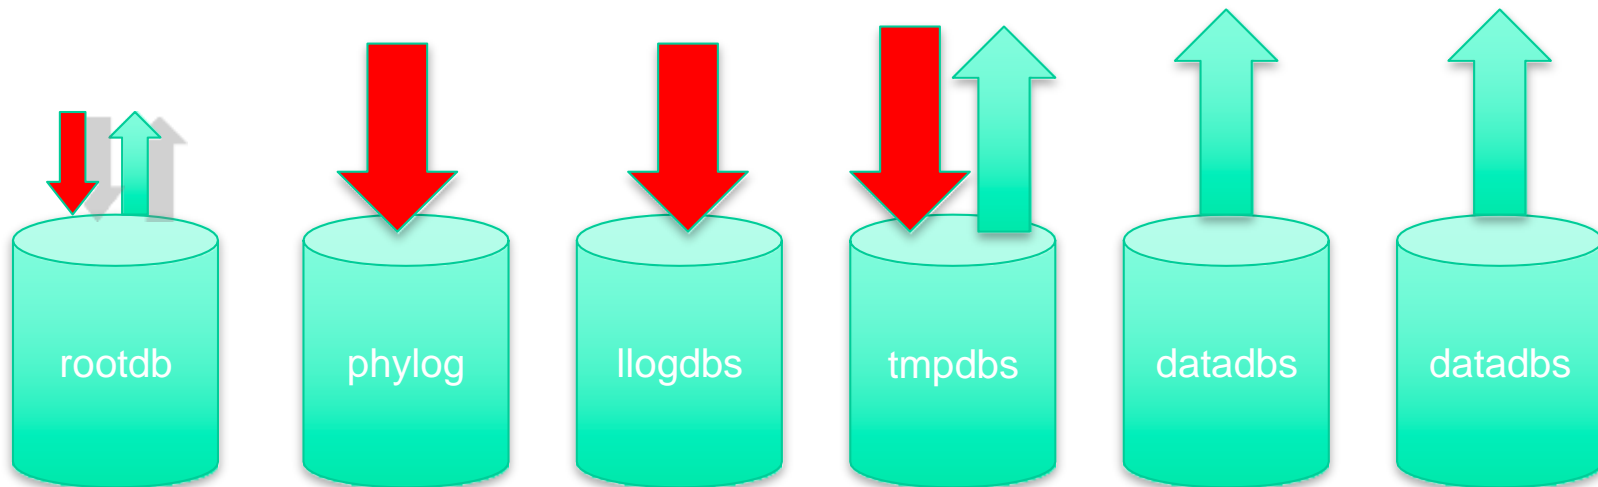
- Create empty file with touch
 - touch \$INFORMIXCHUNKS1/datadbs
- Change the permissions for Informix
 - chmod 660 \$INFORMIXCHUNKS1/datadbs
- Change the owner to informix and group to informix
 - chown informix:informix \$INFORMIXCHUNKS1/datadbs
- Create links if necessary
 - ln -s \$INFORMIXCHUNKS1/datadbs \$INFORMIXLINKS/datadbs

Dbspaces, Chunks, and Pages

- Dbspace
- Chunk
- Page
- Extent
- Tablespace
- Partition
- Fragment

Database Disk I/O

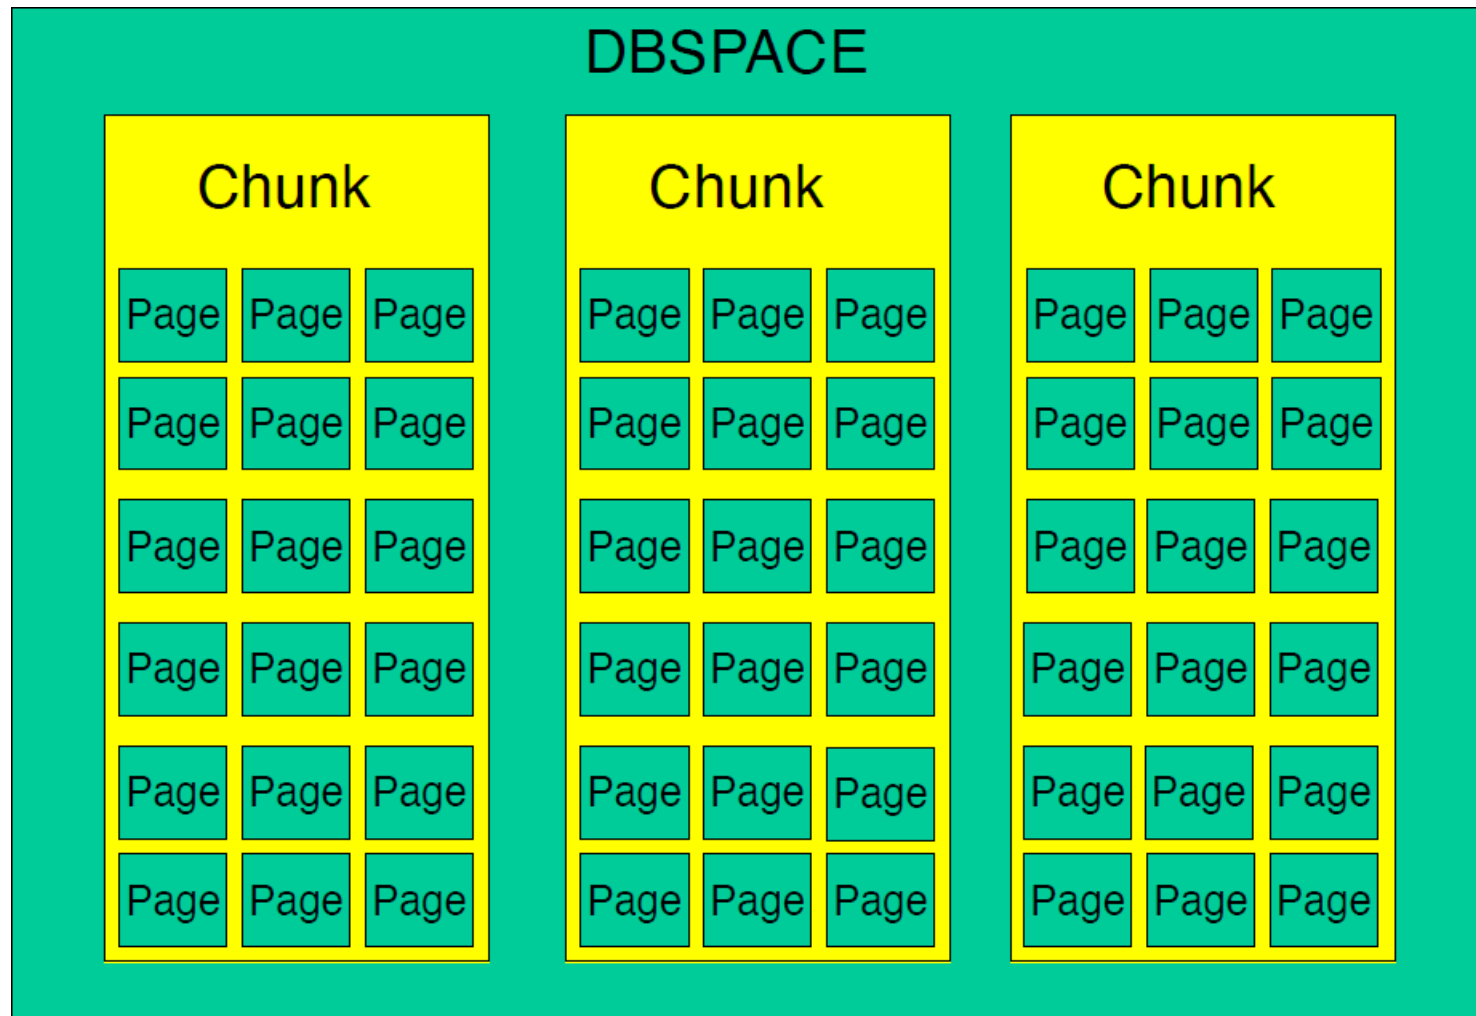
- Most Reads are from Data and Tables
- Writes will be split between Physical Log, Logical Log, Temp, and Data



Managing Disk Space

- **Page** — A page is the physical unit of disk storage that Informix uses to read from and write to Informix databases. The size of a page varies from computer to computer and from dbspace to dbspace. The default page holds either 2 or 4 kilobytes. Because the rootdb, logical log and physical log dbspaces must use the default page size. Since Informix v10.00 the page size for other dbspaces is configurable in multiples of the server's default page size from the default up to 16K.
- **Chunk** — Chunk is the unit of physical disk dedicated to Informix data storage. It represents an allocation of cooked file, cooked device, or raw disk space and is the only unit of physical storage that the Informix administrator allocates.
- **DBspace** — A dbspace is a logical container of one or more chunks.

Disk Layout - Pages, Chunks, and Dbspaces



Page Layout

Page
Header

Slots for rows and
data (HEX 100 or 256
Slots)

Dbospace Considerations

- What is the optimal page size for my data? 2K, 4K, 8K, or 16K?
- Partition table extent sizing (onspaces - ef & -en)
- Should the dbospace be expandable (automate adding chunks from the storage pool)?

Types of DbSPACE

- “Normal” dbspaces
- Temporary dbspaces
- Blobspaces
- Unlogged Smart Blobspaces
- Logged Smart Blobspaces
- Physical Log DbSPACE

Chunk Considerations

- Should the chunk be extendable?
- How active will the data be?
- What else is stored on the device?
- Where should it be placed within storage?
 - SSD
 - Fast magnetic disk
 - Slower magnetic disk
- Should the chunk have an Informix mirror chunk?

Root DBspace

- Define in your ONCONFIG File using the base page size of your Informix port
- Should be on your fastest storage
- Nothing should be in your Root DBspace except
 - Reserved Pages
 - Sysmaster Database
 - Sysutilites Database
 - Sysusers Database

Root DBspace

- Move out of the Root DBSpace
 - Physical Logs
 - Logical Logs
 - Sysadmin Database
 - Temp Dbspace

Physical Log DBspace

- The Physical and Logical log will have 30-50% of all writes
- Move out of Root to separate Dbspaces
- Physical Log Size = $1.25 \times \text{Buffer Size}$
- A Checkpoint will occur when the Physical Log is 75% Full

Logical Log DBspace

- The Physical and Logical log will have 30-50% of all writes
- Move out of Root to separate Dbspaces
- Logical Log Size = Hold 5-10 minutes of transactions at peak time
- Have enough Logical Logs for 4 days

Temp DBspace

- Up to 50% of all I/O can be to the Temp DBspaces (reads + writes)
- Create at least 3 non-logged Temp DBspaces (min) – Informix will use them in round robin and in parallel
- Create 1-3 Logged Temp DBspaces for Logged Temp Transactions if needed

Index DBspace

- Indexes are optimal on 16K page Dbspaces
- Create 16K page size DBSpace for indexes
- Create a 16K BUFFERPOOL for indexes

Data DBspace

- Create Data Dbspaces based on the row size of your tables.
- How many rows will fit on a page?
- What is the least amount of wasted space?
- Spread your data across multiple DBspaces.

How to calculate Optimal Dbospace Size for a Table?

Table Name	Row Size	Rows 2 K page	Waste on 2 K page	Rows 4 K page	Waste on 4 K page	Rows 8 K page	Waste on 8 K page	Rows 16 K page	Waste on 16 K page

- DBspace Page Size (2 K, 4 K, 8 K, 16 K) – 28 bytes
= Size for Data (2048-28=2020)
- Row Size / Dbospace Size for Data = Number of Rows per page
- Waste is how much space is left over
- See Art Kagel's script named "waste" or Lester's Sysmaster script

Partition Large Tables

- Any table larger than a typical chunk may benefit by being partitioned into multiple partitions or tablespaces
- Tables with more than 16,777,216 pages must be partitioned into multiple partitions or fragments
- Partition tables for performance

Informix Storage Pool

- The storage pool contains:
 - the directories, cooked files, and raw devices
 - May be used to automatically expand an existing dbspace
- Storage space threshold is defined in the SP_THRESHOLD parameter
- Database scheduler will automatically run a task that expands the space, either by extending an existing chunk in the space or by adding a new chunk
- Configured in sysadmin database, storagepool table

Informix Storage Pool

- Large Predefined Files to be used to automate expanding DBspaces as needed.
- Current configuration allows only a single storage pool
- Disadvantage: Some applications may require multiple storage pools to allow isolating DBspaces at the storage level. Example:
 - Transaction tables versus historical tables
 - Data tables versus indexes versus blob space versus temp space

Creating an Informix Storage Pool

- Use the Sysadmin task or admin function to create a Storage Pool
 - EXECUTE FUNCTION task("storagepool add", "path", "begin_offset", "total_size", "chunk size", "priority");
- Specify the following:
 - The path for the file, directory, or device to use when additional storage space is required
 - The offset in KB into the device to begin allocating space (0 for cooked files)
 - The total space available to Informix for this device
 - The minimum size in KB of a chunk that can be allocated from the device (1000 KB or greater)
 - The priority of the device (1 = high; 2 = medium; 3 = low) for space allocation

Extendable Chunks

- Since version 11.70 any chunk can be marked “extendable”. When such a chunk fills, the engine will automatically extend the file size a specified amount.
- Only valid for cooked filesystem chunks
- Attribute set after chunk creation with an API function:
`EXECUTE FUNCTION sysadmin:task("modify chunk extendable", "4");`
- Extension size is set by DBspace in an API function:
`EXECUTE FUNCTION task("modify space sp_sizes", "DBspace", "new_sz_kb", "min_extend_sz", "max_size");`

Sets the initial size of new chunks created by dbspace expansion, the maximum size of a chunk for this dbspace, and the amount to extend extendable chunks in this dbspace.

Extendable Chunks Script

```
#####
## Module: @(#)10extendablechunks.sh      2.0      Date: 02/01/2020
## Author: Lester Knutsen  Email: lester@advancedatools.com
##        Advanced DataTools Corporation
##        Description:  Mark Chunks as Extendable
#####
## Setup Environment
echo "Setting up Environment"
. ./informix.env

echo "Making Dbspaces automatically expandable"

dbaccess sysmaster - <<EOF

-- Script to generate SQL to mark the chunks as extendable

output to extendablechunks.sql
without headings
select "execute function sysadmin:task ( 'modify chunk extendable', " || chknum || ");"
from syschunks
    -- Select the dbspaces to make expandable - exclude the following
where dbsnum in
    ( select dbsnum from sysdbspaces where name not in
      ( "rootdbs", "plogdbs", "log1dbs", "log2dbs", "tmp1dbs", "tmp2dbs", "tmp3dbs", "tmp4dbs" )
    );

EOF

dbaccess -e sysadmin extendablechunks.sql
```

Simplify Disk Management

- Dbspaces can be configured to expand by adding new chunks when the existing chunks have been used up.
- Create a storage pool of disk files and devices to use to expand dbspaces.
- Mark dbspaces as expandable.
- If no chunks in those dbspaces are extendable or are extendable but have reached maximum size, the engine will allocate a new chunk from the storage pool.

Disk Layout Best Practices

- Use Symbolic Links
- Avoid Disk IO Contention
- Mirror Critical Media
- Move the Logical and Physical Logs
- Create 3 or more Temp DBspaces
- Isolate High-Use Tables
- Separate Indexes from Data

Use Symbolic Links

- Once a device has been associated with a chunk in Informix it cannot be easily changed (a server restore is required).
- This makes it very difficult to change device names if you need to change disk drives.
- **Recommendation:** use the UNIX facility to create symbolic links and use ONLY symbolic links as the chunk paths passed to onspaces.

Example:

`ln -s /dev/rds1s1 /usr/informix/dev/server1/chunk1`

If the physical disk `/dev/rds1s1` needs to be replaced with a new disk, it may not have the same device name. By using symbolic links, you can create the link to the new device and easily copy the data from `/dev/rds1s1` (using `dd`) or restore Informix from archive.

Avoid Disk IO Contention

- Goals for efficient disk layout typical in a production environment:
 - Limiting disk head movement
 - Reducing disk contention
 - Balancing the load
 - Maximizing availability
- You must make some trade-offs between these goals when you design your disk layout. For example, separating the system catalog tables, the logical logs, and the physical log on physically separate devices can help reduce contention for these resources; however, this action can also increase the chances that you have to perform a system restore.

Mirror Critical Media

- Mirror the critical media – If your chunks are not built from mirrored disk pairs, the root DBspace, the DBspace containing the physical log, and the DBspace containing the logical log files should be mirrored using Informix mirroring. You must specify mirroring on a chunk-by-chunk basis. Locate the primary and the mirrored chunk on different disks. Ideally, different controllers should handle the different disks.
- Mirror speeds up Disk Reads – Informix will read different sectors from both the primary and the mirror at the same time, increasing throughput.

Move the Logical and Physical Logs from the Rootdbs

- The logical log and physical log both contain data that Informix writes frequently. Likewise, reserved pages are read frequently; they contain internal tables that describe and track all DBspaces, blobspaces, chunks, databases, and tblspaces.
- By default, Informix stores the logical and physical logs together in the root DBspace when a new server is initialized. Maintaining these together in the root DBspace will become a source of contention as your database system grows.
- Reduce this contention and provide better load balancing by moving the logical and physical logs to separate partitions or, even better, separate disk drives. For optimum performance, create two or more additional DBspaces:
 - one for the physical log (take advantage of the new physical log DBspace type v12.10+)
 - one or more for the logical log (if more than one alternate creating logs round robin).
- When you move the logs, avoid storing them in a db space/disk that contains high-access rate tables; instead consider storing them in a DBspace dedicated to storing only the physical or logical log.

Create 3 or More Temp DBspaces

- Informix will read/write multiple temp DBspaces in parallel and create temp tables fragmented across them.
- Sort-work files are written round robin to all temp DBspaces.
- Merging sort-work files will read from two temp DBspaces and write to a third if available.
- Temp DBspaces may have high levels of activity
- Move temp DBspaces to separate disks
- ONCONFIG file

```
DBSPACETEMP tmp1dbs:tmp2dbs:tmp3dbs:normal_dbspc
```

Isolate High-Use Tables

- Place a table with high I/O activity on a disk device dedicated to its use and thus reduce contention.
- Put the tables partitions with the highest frequency of use on the fastest drives.
- Placing two high-access tables on separate disk devices reduces competition for disk access when joins are formed between the two tables or when the two tables experience frequent, simultaneous access from multiple applications.

Separate Indexes from Data

- Create a DBspace for data
- Create a separate DBspace for indexes
- Informix will read indexes and data in parallel

Monitoring Disk Performance

Onstat Command	Description
onstat -d	Print dbspaces and chunks
onstat -D	Print dbspaces and chunk IO
onstat -g iof	Print disk IO statistics by chunk/file
onstat -g iov	Print disk IO statistics by vp
onstat -g ioh	Print IO history for the last hour by chunk

Monitor with Onstat -d

```
informix@tiger1:~/InformixAdvclass/lab09-extra train1 > onstat -d
```

```
IBM Informix Dynamic Server Version 14.10.FC3 -- On-Line -- Up 00:16:38 -- 3606768 Kbytes
```

Dbspaces

address	number	flags	fchunk	nchunks	pgsize	flags	owner	name
4a651028	1	0x4020001	1	1	2048	N BA	informix	rootdbs
4ceeade8	2	0x4020001	2	1	2048	N BA	informix	logdbs
4cb5ed98	3	0x4020001	3	1	2048	N BA	informix	datadbs
4c735508	4	0x4002001	4	1	2048	N TBA	informix	tmpdbs
4c6aad38	5	0x4020001	5	1	2048	N BA	informix	datab3dbs

5 active, 2047 maximum

Chunks

address	chunk/dbs	offset	size	free	bpages	flags	pathname
4a651268	1 1	0	1000000	739855		P0-B--	/informixchunks/train1/rootdbs
4be12028	2 2	0	1000000	199947		P0-B--	/informixchunks/train1/logdbs
4ce9f028	3 3	0	2000000	982726		P0-B--	/informixchunks/train1/datadbs
4ce9a028	4 4	0	1000000	999947		P0-B--	/informixchunks/train1/tmpdbs
4c7e3028	5 5	0	5000000	2647971		P0-B--	/informixchunks/train1/datab3dbs

5 active, 32766 maximum

NOTE: The values in the "size" and "free" columns for DBspace chunks are displayed in terms of "pgsize" of the DBspace to which they belong.

Expanded chunk capacity mode: always

Monitor with Onstat -D

```
informix@tiger1:~/InformixAdvclass/lab09-extra train1 > onstat -D
```

```
IBM Informix Dynamic Server Version 14.10.FC3 -- On-Line -- Up 00:19:07 -- 3606768 Kbytes
```

Dbspaces

address	number	flags	fchunk	nchunks	pgsize	flags	owner	name
4a651028	1	0x4020001	1	1	2048	N BA	informix	rootdbs
4ceeade8	2	0x4020001	2	1	2048	N BA	informix	logdbs
4cb5ed98	3	0x4020001	3	1	2048	N BA	informix	datadbs
4c735508	4	0x4002001	4	1	2048	N TBA	informix	tmpdbs
4c6aad38	5	0x4020001	5	1	2048	N BA	informix	datab3dbs

5 active, 2047 maximum

Chunks

address	chunk/dbs	offset	page Rd	page Wr	pathname
4a651268	1 1	0	1504	304360	/informixchunks/train1/rootdbs
4be12028	2 2	0	0	1631495	/informixchunks/train1/logdbs
4ce9f028	3 3	0	765962	1085587	/informixchunks/train1/datadbs
4ce9a028	4 4	0	0	4	/informixchunks/train1/tmpdbs
4c7e3028	5 5	0	11375662	2264392	/informixchunks/train1/datab3dbs

5 active, 32766 maximum

Is your IO Spread out on all chunks?

NOTE: The values in the "page Rd" and "page Wr" columns for DSpace chunks are displayed in terms of system base page size.

Monitor with onstat -g iof

IBM Informix Dynamic Server Version 14.10.FC3 -- On-Line -- Up 00:19:57 -- 3606768 Kbyte

AIO global files:

gfd	pathname	bytes read	page reads	bytes write	page writes	io/s
3	rootdbs	3080192	1504	623329280	304360	719.7
	op type	count	avg. time			
	seeks	0	N/A			
	reads	1466	0.0000			
	writes	6612	0.0017			
	kaio_reads	0	N/A			
	kaio_writes	0	N/A			
4	logdbs	0	0	3341301760	1631495	2462.7
	op type	count	avg. time			
	seeks	0	N/A			
	reads	0	N/A			
	writes	28402	0.0004			
	kaio_reads	0	N/A			
	kaio_writes	0	N/A			
5	datadbs	1568690176	765962	2223282176	1085587	72124.2
	op type	count	avg. time			
	seeks	0	N/A			
	reads	761698	0.0000			
	writes	68087	0.0001			
	kaio_reads	0	N/A			
	kaio_writes	0	N/A			
6	tmpdbs	0	0	8192	4	188786.3
	op type	count	avg. time			
	seeks	0	N/A			
	reads	0	N/A			
	writes	2	0.0000			
	kaio_reads	0	N/A			
	kaio_writes	0	N/A			

Shows no KAIO

Monitor > 0.01

Monitor with onstat -g iov

IBM Informix Dynamic Server Version 14.10.FC3 -- On-Line -- Up 00:23:08 -- 3606768 Kbytes

AIO I/O vps:

class	vp	id	s	io/s	totalops	dskread	dskwrite	dskcopy	wakeups	io/wup	errors	tempops
fifo	7	0	i	0.0	0	0	0	0	1	0.0	0	0
msc	6	0	i	0.0	7	0	0	0	8	0.9	0	7
aio	5	0	i	1140.8	1583436	853572	729079	0	1032932	1.5	0	0
aio	12	1	i	397.9	552218	4422	547782	0	5357	103.1	0	0
pio	4	0	i	0.1	120	0	120	0	121	1.0	0	120
lio	3	0	i	22.3	30911	0	30911	0	30912	1.0	0	30911



Is the IO balance among AIO VPs?

Monitor with onstat -g ioh

IBM Informix Dynamic Server Version 14.10.FC3 -- On-Line -- Up 00:24:40 -- 3606768 Kbytes

AIO global files:

gfd	pathname	bytes read	page reads	bytes write	page writes	io/s
3	rootdbs	3080192	1504	623329280	304360	719.8

time	avg read			avg write		
	reads	io/s	op time	writes	io/s	op time
12:46:04	0	0.0	0.00000	0	0.0	0.00000
12:45:04	0	0.0	0.00000	0	0.0	0.00000
12:44:04	0	0.0	0.00000	0	0.0	0.00000
12:43:04	0	0.0	0.00000	0	0.0	0.00000
12:42:04	0	0.0	0.00000	0	0.0	0.00000
12:41:04	3	0.1	0.00001	15	0.2	0.00024
12:40:04	0	0.0	0.00000	0	0.0	0.00000
12:39:04	0	0.0	0.00000	0	0.0	0.00000
12:38:04	1	0.0	0.00001	0	0.0	0.00000
12:37:04	0	0.0	0.00000	0	0.0	0.00000
12:36:04	15	0.2	0.00001	5	0.1	0.00012
12:35:04	0	0.0	0.00000	0	0.0	0.00000
12:34:04	0	0.0	0.00000	0	0.0	0.00000
12:33:04	0	0.0	0.00000	0	0.0	0.00000
12:32:04	0	0.0	0.00000	0	0.0	0.00000
12:31:04	24	0.4	0.00001	16	0.3	0.00011
12:30:04	16	0.3	0.00001	2	0.0	0.00014
12:29:04	2	0.0	0.00000	0	0.0	0.00000
12:28:04	0	0.0	0.00000	15	0.2	0.00005
12:27:04	0	0.0	0.00000	0	0.0	0.00000
12:26:04	9	0.1	0.00001	31	0.5	0.00070
12:25:04	0	0.0	0.00000	44	0.7	0.00007
12:24:04	8	0.1	0.00001	163	2.7	0.00009
12:23:04	1388	23.1	0.00000	6321	105.3	0.00144

Monitor Unix Disk IO with iostat

```
informix@tiger1:~/InformixAdvclass/lab09-extra train1 > iostat 5 5
Linux 3.10.0-1062.12.1.el7.x86_64 (tiger1)      03/16/2020      _x86_64_
```

avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle
	0.03	0.00	0.02	0.01	0.00	99.94

Device:	tps	kB_read/s	kB_wrtn/s	kB_read	kB_wrtn
sda	1.88	5.38	288.48	9813895	525761373
dm-0	0.11	0.66	0.72	1210035	1315866
dm-1	0.00	0.00	0.00	2204	1628
dm-2	0.02	2.07	0.05	3767548	86698
dm-3	1.81	2.64	287.70	4817011	524340608

avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle
	0.08	0.00	0.03	0.03	0.00	99.87

Device:	tps	kB_read/s	kB_wrtn/s	kB_read	kB_wrtn
sda	1.40	5.60	3.20	28	16
dm-0	0.40	4.80	0.00	24	0
dm-1	0.00	0.00	0.00	0	0
dm-2	0.00	0.00	0.00	0	0
dm-3	1.00	0.80	3.20	4	16

The onspaces Command

ONSPACES

Usage: onspaces -a <spacename> -p <path> -o <offset> -s <size> [-m <path> <offset>]
 { { [-Mo <mdoffset>] [-Ms <mdsize>] } | -U } |
 -c -d <DBspace> [-k <pagesize>] [-t] -p <path> -o <offset> -s <size>
 [-m <path> <offset>] |
 -c -d <DBspace> [-k <pagesize>] -p <path> -o <offset> -s <size>
 [-m <path> <offset>] [-ef <first_extent_size>] [-en <next_extent_size>] |
 -c -b <BLOBspace> -g <pagesize> -p <path> -o <offset> -s <size>
 [-m <path> <offset>] |
 -c -P <PLOGspace> -p <path> -o <offset> -s <size> [-m <path> <offset>] |
 -c -S <SBLOBspace> [-t] -p <path> -o <offset> -s <size> [-m <path> <offset>]
 [-Mo <mdoffset>] [-Ms <mdsize>] [-Df <default-list>] |
 -c -x <Extspace> -l <Location> [-d <spacename>] [-p <path> -o <offset>] [-f] [-y] |
 -f[y] off [<DBspace-list>] | on [<DBspace-list>] |
 -m <spacename> {-p <path> -o <offset> -m <path> <offset> [-y] | -f <filename>} |
 -r <spacename> [-y] |
 -s <spacename> -p <path> -o <offset> {-O | -D} [-y] |
 -ch <sbspacename> -Df <default-list> |
 -cl <sbspacename> |
 -ren <spacename> -n <newname>
 -a Add a chunk to a DBspace, BLOBspace or SBLOBspace
 -c Create a DBspace, BLOBspace, SBLOBspace or Extspace
 -d Drop a DBspace, BLOBspace, SBLOBspace, Extspace, or chunk
 -f Change dataskip default for specified DBspaces
 -m Add mirroring to an existing DBspace, BLOBspace or SBLOBspace
 -r Turn mirroring off for a DBspace, BLOBspace or SBLOBspace
 -s Change the status of a chunk
 -ch Change default list for smart large object space
 -cl garbage collect smart large objects that are not referenced default-list = {[LOGGING =
 {ON|OFF}] [,ACcesstime = {ON|OFF}] [,AVG_LO_SIZE = {1 - 2097152}]}
 -ren Rename a DBspace, BLOBspace, SBLOBspace or Extspace
 -u Create the new space unencrypted

Using Sysadmin Task

```
-- Dbspace 2 -- Chunk 2
EXECUTE FUNCTION TASK
('create plogspace', 'plogdbs', '/informixchunks/newserver2/informixlinks/plogdbs', '4000000', '0');

-- Dbspace 3 -- Chunk 3
EXECUTE FUNCTION TASK
('create dbspace', 'log1dbs', '/informixchunks/newserver2/informixlinks/log1dbs', '2000000', '0', '2', '100', '100');

-- Dbspace 4 -- Chunk 4
EXECUTE FUNCTION TASK
('create dbspace', 'log2dbs', '/informixchunks/newserver2/informixlinks/log2dbs', '2000000', '0', '2', '100', '100');

-- Dbspace 5 -- Chunk 5
EXECUTE FUNCTION TASK
('create tempdbspace', 'tmp1dbs', '/informixchunks/newserver2/informixlinks/tmp1dbs', '2000000', '0', '2', '100', '100');

-- Dbspace 6 -- Chunk 6
EXECUTE FUNCTION TASK
('create tempdbspace', 'tmp2dbs', '/informixchunks/newserver2/informixlinks/tmp2dbs', '2000000', '0', '2', '100', '100');

-- Dbspace 7 -- Chunk 7
EXECUTE FUNCTION TASK
('create tempdbspace', 'tmp3dbs', '/informixchunks/newserver2/informixlinks/tmp3dbs', '2000000', '0', '2', '100', '100');

-- Dbspace 8 -- Chunk 8
EXECUTE FUNCTION TASK
('create tempdbspace', 'tmp4dbs', '/informixchunks/newserver2/informixlinks/tmp4dbs', '2000000', '0', '2', '100', '100');

-- Dbspace 9 -- Chunk 9
EXECUTE FUNCTION TASK
('create dbspace', 'sysadmdbs', '/informixchunks/newserver2/informixlinks/sysadmdbs', '2000000', '0', '2', '100', '400');
```

Using onspaces

```
# Dbspace 2 -- Chunk 2
onspaces -c -P plogdbs -p /informixchunks/newserver2/informixlinks/plogdbs -o 0 -s 4000000

# Dbspace 3 -- Chunk 3
onspaces -c -d log1dbs -k 2 -p /informixchunks/newserver2/informixlinks/log1dbs -o 0 -s 2000000 -ef 100 -en 100

# Dbspace 4 -- Chunk 4
onspaces -c -d log2dbs -k 2 -p /informixchunks/newserver2/informixlinks/log2dbs -o 0 -s 2000000 -ef 100 -en 100

# Dbspace 5 -- Chunk 5
onspaces -c -d tmp1dbs -k 2 -t -p /informixchunks/newserver2/informixlinks/tmp1dbs -o 0 -s 2000000

# Dbspace 6 -- Chunk 6
onspaces -c -d tmp2dbs -k 2 -t -p /informixchunks/newserver2/informixlinks/tmp2dbs -o 0 -s 2000000

# Dbspace 7 -- Chunk 7
onspaces -c -d tmp3dbs -k 2 -t -p /informixchunks/newserver2/informixlinks/tmp3dbs -o 0 -s 2000000

# Dbspace 8 -- Chunk 8
onspaces -c -d tmp4dbs -k 2 -t -p /informixchunks/newserver2/informixlinks/tmp4dbs -o 0 -s 2000000

# Dbspace 9 -- Chunk 9
onspaces -c -d sysadmdbs -k 2 -p /informixchunks/newserver2/informixlinks/sysadmdbs -o 0 -s 2000000 -ef 100 -en 400

# Dbspace 10 -- Chunk 10
onspaces -c -d datadbs -k 2 -p /informixchunks/newserver2/informixlinks/datadbs -o 0 -s 2000000 -ef 100 -en 100

# Dbspace 11 -- Chunk 11
onspaces -c -d data1dbs -k 2 -p /informixchunks/newserver2/informixlinks/data1dbs -o 0 -s 2000000 -ef 100 -en 100

# Dbspace 12 -- Chunk 12
onspaces -c -d data2dbs -k 2 -p /informixchunks/newserver2/informixlinks/data2dbs -o 0 -s 2000000 -ef 100 -en 100
```